

Rice Type Classification Using CNN

Date	20 July 2024
Team ID	SWTID1720110358
Project Name	Rice Type Classification using CNN

1. Introduction

Rice is a staple food consumed globally, with various types exhibiting unique physical and chemical properties. Accurate classification of rice types is crucial for quality control, pricing, and meeting consumer preferences. Traditional methods of rice classification are labour-intensive and prone to errors. This project leverages Convolutional Neural Networks (CNN) to automate the classification process, providing a fast, accurate, and scalable solution.

1.1 Project overview

1. Data Collection

- Collect images of various rice types from multiple sources.
- Ensure accurate labeling of the images according to rice type.

2. Data Preprocessing

- Resize images to a consistent dimension.
- Normalize pixel values for uniform input to the model.
- Perform data augmentation to improve model generalization (e.g., rotations, flips).

3. Model Development

- Design a CNN architecture suitable for image classification.
- Consider using well-known architectures (e.g., VGG16, ResNet) or developing a custom model.

4. Model Training

- Split the dataset into training, validation, and test sets.
- Train the CNN model using the training set and validate its performance on the validation set.
- Fine-tune hyperparameters to optimize model performance.

5. Model Evaluation

- Test the model on the test set.
- Use evaluation metrics such as accuracy, precision, recall, and F1 score to assess performance.

6. Deployment

- Develop a user-friendly interface for model interaction.

- Deploy the model in a real-world setting for automated rice classification.

4. Tools and Technologies

- **Programming Languages:** Python
- **Libraries and Frameworks:** TensorFlow/Keras or PyTorch, OpenCV
- **Development Environment:** Jupyter Notebook or any suitable Python IDE
- **Hardware:** GPU (optional but recommended for faster training)

5. Potential Challenges

- **Data Quality:** Ensuring high-quality, well-labeled images.
- **Model Generalization:** Preventing overfitting to maintain performance on new data.

1.2 Objectives

- **Develop a CNN model** capable of classifying different types of rice based on images.
- **Create a labelled dataset** of rice images encompassing various types.
- **Evaluate the performance** of the CNN model using metrics like accuracy, precision, recall, and F1 score.
- **Deploy the model** for practical use in quality control processes.

Project Initialization and Planning Phase

Date	15 March 2024
Team ID	SWTID1720110358
Project Name	Rice Classification using CNN
Maximum Marks	3 Marks

Define Problem Statements (Customer Problem Statement Template):

There are many types of rice available for production. It is essential to identify the type of rice as each produce needs different amounts of water, manure, etc.

It is not possible for the farmers to pay the agriculture experts hefty fees every time they have a new produce. We have come up with a solution to this problem. We have trained an AI model which can be used by farmers to check the type of rice. The users need to upload image of a rice grain and click on the submit button. Our model will give its prediction for probable rice type based on the image. Our model can predict up to 5 different types of rice.

This model is useful for farmers, agriculture scientists, home farmers, gardeners, etc. This AI model is made using Convolutional Neural networks and under CNN we will be using transfer learning. Transfer learning has become one of the most common techniques that has achieved better performance in many areas, especially in image analysis and classification. We used Transfer Learning technique MobileNetv4 that is more widely used as a transfer learning method in image analysis, and it is highly effective.

Problem Statement (PS)	I am (Customer)	I'm trying to	But	Because	Which makes me feel
PS-1	Dhruv Garg	Find the type of rice I have	I am not sure the type of rice I have	There are different shapes and types of rice	Confused

Project Initialization and Planning Phase

Date	15 March 2024
Team ID	SWTID1720110358
Project Title	Rice Classification using CNN
Maximum Marks	3 Marks

Project Proposal (Proposed Solution):

We have come up with a solution to this problem. We have trained an AI model which can be used by farmers to check the type of rice. The users need to upload image of a rice grain and click on the submit button. Our model will give its prediction for probable rice type based on the image. Our model can predict up to 5 different types of rice.

Project Overview	
Objective	To detect the types of rice given by the user
Scope	This model is useful for farmers, agriculture scientists, home farmers, gardeners, etc. This AI model is made using Convolutional Neural networks and under CNN we will be using transfer learning.
Problem Statement	
Description	This AI model is made using Convolutional Neural networks and under CNN we will be using transfer learning. Transfer learning has become one of the most common techniques that has achieved better performance in many areas, especially in image analysis and classification.
Impact	it is essential to identify the type of rice as each produce needs different amounts of water, manure, etc. It is not possible for the farmers to pay the agriculture experts hefty fees every time they have a new produce.
Proposed Solution	
Approach	Transfer learning has become one of the most common techniques that has achieved better performance in many areas, especially in image analysis and classification. We used Transfer Learning

	technique MobileNetv4 that is more widely used as a transfer learning method in image analysis, and it is highly effective.
Key Features	<ul style="list-style-type: none"> The MobileNet Model analyzes the image, then the prediction is showcased on the Flask UI.

Resource Requirements

Resource Type	Description	Specification/Allocation
Hardware		
Computing Resources	CPU/GPU specifications, number of cores	e.g., 2 x NVIDIA V100 GPUs
Memory	RAM specifications	e.g., 8 GB
Storage	Disk space for data, models, and logs	e.g., 1 TB SSD
Software		
Frameworks	Python frameworks	e.g., Flask
Libraries	Additional libraries	e.g., tensorflow
Development Environment	IDE, version control	e.g., Jupyter Notebook, Git
Data		
Data	Source, size, format	e.g., Kaggle dataset, 10,000 images

Data Collection and Preprocessing Phase

Date	15 March 2024
Team ID	SWTID1720110358
Project Title	Rice Classification using CNN
Maximum Marks	2 Marks

Data Collection Plan & Raw Data Sources Identification Template

Elevate your data strategy with the Data Collection plan and the Raw Data Sources report, ensuring meticulous data curation and integrity for informed decision-making in every analysis and decision-making endeavor.

Data Collection Plan Template

Section	Description
Project Overview	A machine learning project utilizing Convolutional Neural Networks (CNN) to classify different types of rice grains accurately, aiming to improve efficiency and accuracy in rice quality assessment and sorting processes
Data Collection Plan	Kaggle https://www.muratkoklu.com/datasets/
Raw Data Sources Identified	Description: This dataset contains labeled images of five different types of rice grains. Types of Rice: Arborio, Basmati, Jasmine, Ipsala, and Karacadag.

	<p>Content: High-resolution images, each labeled with the corresponding type of rice.</p> <p>Usage: Ideal for training and testing Convolutional Neural Networks (CNN) for rice classification tasks.</p>
--	---

Raw Data Sources Template

Source Name	Description	Location/URL	Format	Size	Access Permissions
Dataset 1	This dataset has 75K images including 15K pieces from each rice variety.	https://www.muratkoklu.com/datasets/	Image	0.24 GB	Public
Dataset 2	This dataset has 12 morphological, 4 shape and 90 color features.	https://www.muratkoklu.com/datasets/	Image	0.24 GB	Public

Data Collection and Preprocessing Phase

Date	15 March 2024
Team ID	SWTID1720110358
Project Title	Rice Classification using CNN
Maximum Marks	2 Marks

Data Quality Report Template

The Data Quality Report Template will summarize data quality issues from the selected source, including severity levels and resolution plans. It will aid in systematically identifying and rectifying data discrepancies.

Data Source	Data Quality Issue	Severity	Resolution Plan
Dataset	Inconsistent image sizes	Moderate	Resize all images to a uniform size using image processing libraries (e.g., OpenCV or PIL) before feeding them into the CNN.
Dataset	Variations in lighting and shadows	Moderate	Use data augmentation techniques (e.g., brightness adjustment, rotation) to create a more robust model.

Data Collection and Preprocessing Phase

Date	15 March 2024
Team ID	SWTID1720110358
Project Title	Rice Classification using CNN
Maximum Marks	6 Marks

Preprocessing Template

The images will be preprocessed by resizing, normalizing, augmenting, denoising, adjusting contrast, detecting edges, converting color space, cropping, batch normalizing, and whitening data. These steps will enhance data quality, promote model generalization, and improve convergence during neural network training, ensuring robust and efficient performance across various computer vision tasks.

Section	Description
Data Overview	Images of different types of rice
Resizing	Resize images to a specified target size.
Normalization	Normalize pixel values to a specific range.
Data Augmentation	Apply augmentation techniques such as flipping, rotation, shifting, zooming, or shearing.

Data Preprocessing Code Screenshots

Loading Data	<pre>In [3]: data_dir=Path(r'E:\DataScience\Rice Type Detection\RiceClassification\Data\Rice_Image_Dataset') arborio=list(data_dir.glob('Arborio/*'))[:600] basmati=list(data_dir.glob('Basmati/*'))[:600] ipsala=list(data_dir.glob('Ipsala/*'))[:600] jasmine=list(data_dir.glob('Jasmine/*'))[:600] karacadag=list(data_dir.glob('Karacadag/*'))[:600]</pre>
--------------	---

<h2>Resizing</h2>	<div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> </div> <pre> In [6]: X, y = [], [] for label, images in df_images.items(): for image in images: img=cv2.imread(str(image)) resized_img=cv2.resize(img, (224,224)) X.append(resized_img) y.append(df_labels[label]) In [7]: X=np.array(X) X=X/255 y=np.array(y) </pre>
<h2>Normalization</h2>	<div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> </div> <pre> In [6]: X, y = [], [] for label, images in df_images.items(): for image in images: img=cv2.imread(str(image)) resized_img=cv2.resize(img, (224,224)) X.append(resized_img) y.append(df_labels[label]) In [7]: X=np.array(X) X=X/255 y=np.array(y) </pre>
<h2>Data Augmentation</h2>	<div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> <div>0 50 100 150 200</div> </div> <pre> In [6]: X, y = [], [] for label, images in df_images.items(): for image in images: img=cv2.imread(str(image)) resized_img=cv2.resize(img, (224,224)) X.append(resized_img) y.append(df_labels[label]) In [7]: X=np.array(X) X=X/255 y=np.array(y) </pre>

Model Development Phase Template

Date	15 March 2024
Team ID	SWTID1720110358
Project Title	Rice Classification using CNN
Maximum Marks	5 Marks

Model Selection Report

MobileNetV2 is selected for the model primarily due to its exceptional balance between efficiency and performance. MobileNetV2 is designed to be efficient in terms of both speed and memory usage. Despite its efficiency, MobileNetV2 maintains high accuracy on various image classification tasks. It strikes a good balance between performance and computational requirements. MobileNetV2 can be easily integrated into a larger model and fine-tuned for specific tasks.

Model Selection Report:

Model	Description
Model 1	A sequential model using MobileNet as the base, with an additional dense layer for classification into 5 labels.

Model Development Phase Template

Date	15 March 2024
Team ID	SWTID1720110358
Project Title	Rice Classification using CNN
Maximum Marks	10 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include a summary and training and validation performance metrics for multiple models, presented through respective screenshots.

Initial Model Training Code (5 marks):

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.applications import MobileNet

num_label = 5 # number of labels

# Load the MobileNet model without the top classification layer
mobile_net = MobileNet(include_top=False, input_shape=(224, 224, 3), pooling='avg')

# Create a Sequential model and add MobileNet as the base
model = models.Sequential([
    mobile_net,
    layers.Dense(num_label, activation='softmax')
])

model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet/mobilenet_1_0_224_tf_no_top.h5
17225924/17225924 ————— 3s 0us/step
odel: "sequential_4"

Layer (type)	Output Shape	Param #
mobilenet_1.00_224 (Functional)	?	3,228,864
dense_4 (Dense)	?	0 (unbuilt)

Total params: 3,228,864 (12.32 MB)

Trainable params: 3,206,976 (12.23 MB)

Non-trainable params: 21,888 (85.50 KB)

```
In [15]: model.compile(
optimizer="adam",
loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
metrics=['acc'])
```

```
In [16]: history = model.fit(X_train, y_train, epochs=10, validation_data=(X_val, y_val))
```

Epoch 1/10

Model Validation and Evaluation Report (5 marks):

Model	Summary	Training and Validation Performance Metrics									
Model 1	<p>Downloading data from https://storage.googleapis.com/tensorflow/tfjs-applications/mobilenet_v1_0.25_v2_top_01.tgz</p> <p>Model: "sequential_4"</p> <table border="1"> <thead> <tr> <th>Layer (type)</th><th>Output Shape</th><th>Param #</th></tr> </thead> <tbody> <tr> <td>mobilenet_v1_0.25_v2_top_01 (Functional)</td><td>?</td><td>1,228,864</td></tr> <tr> <td>dense_4 (Dense)</td><td>?</td><td>0 (unbuilt)</td></tr> </tbody> </table> <p>Total params: 1,228,864 (12.12 MB) Trainable params: 1,228,864 (12.12 MB) Non-trainable params: 0 (0.00 MB)</p>	Layer (type)	Output Shape	Param #	mobilenet_v1_0.25_v2_top_01 (Functional)	?	1,228,864	dense_4 (Dense)	?	0 (unbuilt)	<pre>Epoch 1/10: 216s 3s/step - acc: 0.8866 - loss: 0.3554 - val_acc: 0.7394 - val_loss: 1.1981 Epoch 2/10: 159s 2s/step - acc: 0.9748 - loss: 0.0907 - val_acc: 0.7872 - val_loss: 0.9475 Epoch 3/10: 159s 2s/step - acc: 0.9720 - loss: 0.1120 - val_acc: 0.9574 - val_loss: 0.1104 Epoch 4/10: 159s 2s/step - acc: 0.9938 - loss: 0.0230 - val_acc: 1.0000 - val_loss: 0.0018 Epoch 5/10: 157s 2s/step - acc: 0.9944 - loss: 0.0154 - val_acc: 0.9947 - val_loss: 0.0139 Epoch 6/10: 158s 2s/step - acc: 0.9993 - loss: 0.0035 - val_acc: 0.9521 - val_loss: 0.1182 Epoch 7/10: 163s 2s/step - acc: 0.9919 - loss: 0.0305 - val_acc: 1.0000 - val_loss: 0.0017 Epoch 8/10: 165s 2s/step - acc: 0.9915 - loss: 0.0209 - val_acc: 1.0000 - val_loss: 0.0124 Epoch 9/10: 194s 3s/step - acc: 0.9937 - loss: 0.0203 - val_acc: 0.9947 - val_loss: 0.0098 Epoch 10/10: 194s 3s/step - acc: 0.9985 - loss: 0.0032 - val_acc: 1.0000 - val_loss: 8.0853e-05</pre>
Layer (type)	Output Shape	Param #									
mobilenet_v1_0.25_v2_top_01 (Functional)	?	1,228,864									
dense_4 (Dense)	?	0 (unbuilt)									

Model Optimization and Tuning Phase Template

Date	15 March 2024
Team ID	SWTID1720110358
Project Title	Rice Type Classification using CNN
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining neural network models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (8 Marks):

Model	Tuned Hyperparameters
Model 1	Hyperparam1, Hyperparam2, ... (Give a short description of the params and give the screenshot of the code)
Model 2	Hyperparam1, Hyperparam2, ... (Give a short description of the params and give the screenshot of the code)
...	...

Hyperparameter Tuning Documentation:

Model 1: USING CNN (includes Softmax Regression & Adam Optimiser)

Key Hyperparameters

1. **Learning Rate (`learning_rate`):**
 - Controls the step size during gradient descent.
 - Typical values: 0.00001, 0.0001, 0.001.
2. **Batch Size (`batch_size`):**
 - Number of samples processed before the model is updated.
 - Common values: 16, 32, 64.
3. **Number of Epochs (`epochs`):**
 - Number of complete passes through the training dataset.
 - Usually between 10 to 50 for fine-tuning.
4. **Dropout Rate (`dropout_rate`):**
 - Fraction of the input units to drop for preventing overfitting.
 - Common values: 0.3, 0.4, 0.5.
5. **Optimizer (`optimizer`):**
 - Algorithm used to minimize the loss function.
 - Adam optimizer is adaptive and commonly used.
6. **Base Model Trainability:**
 - Whether to fine-tune the layers of the pre-trained model.
 - Setting `base_model.trainable` to False initially, then to True for fine-tuning.

TESTING THE MODEL :

```
[29] a1 = cv2.imread("E:\DataScience\Rice Type Detection\RiceClassification\Data\Rice_Image_Dataset\Arborio\Arborio (3).jpg")
Python

[30] a1 = cv2.resize(a1,(224,224))
Python

[31] a1 = np.array(a1)
Python

[32] a1 = a1/255
Python

[33] a1 = np.expand_dims(a1, 0)
Python

[33] a1 = np.expand_dims(a1, 0)
Python

[34] pred = model.predict(a1)
Python

... 1/1 — 0s 45ms/step

[36] pred = pred.argmax()
pred
Python

... 0

[37] for i, j in df_labels.items():
    if pred == j:
        print(i)
Python

... arborio
```

ACCURACY:


```
history = model.fit(X_train, y_train, epochs=10, validation_data=(X_val, y_val))
```

Python

Epoch 1/10
c:\Users\Dhruv Garg\anaconda3\lib\site-packages\keras\src\backend\tensorflow\nn.py:625: UserWarning: "`sparse_categorical_crossentropy, from_logits = _get_logits(
71/71 ----- 216s 3s/step - acc: 0.8866 - loss: 0.3554 - val_acc: 0.7394 - val_loss: 1.1981
Epoch 2/10
71/71 ----- 159s 2s/step - acc: 0.9748 - loss: 0.0907 - val_acc: 0.7872 - val_loss: 0.9475
Epoch 3/10
71/71 ----- 159s 2s/step - acc: 0.9720 - loss: 0.1120 - val_acc: 0.9574 - val_loss: 0.1104
Epoch 4/10
71/71 ----- 159s 2s/step - acc: 0.9938 - loss: 0.0230 - val_acc: 1.0000 - val_loss: 0.0018
Epoch 5/10
71/71 ----- 157s 2s/step - acc: 0.9944 - loss: 0.0154 - val_acc: 0.9947 - val_loss: 0.0139
Epoch 6/10
71/71 ----- 158s 2s/step - acc: 0.9993 - loss: 0.0035 - val_acc: 0.9521 - val_loss: 0.1182
Epoch 7/10
71/71 ----- 163s 2s/step - acc: 0.9919 - loss: 0.0305 - val_acc: 1.0000 - val_loss: 0.0017
Epoch 8/10
71/71 ----- 165s 2s/step - acc: 0.9915 - loss: 0.0209 - val_acc: 1.0000 - val_loss: 0.0124
Epoch 9/10
71/71 ----- 194s 3s/step - acc: 0.9937 - loss: 0.0203 - val_acc: 0.9947 - val_loss: 0.0098
Epoch 10/10
71/71 ----- 194s 3s/step - acc: 0.9985 - loss: 0.0032 - val_acc: 1.0000 - val_loss: 8.0853e-05

PRECISION, RECALL, F1 SCORE:

```
model.evaluate(X_test,y_test)
```

Python

18/18 ----- 8s 447ms/step - acc: 1.0000 - loss: 0.0025
[0.0019843829795718193, 1.0]

```
from sklearn.metrics import classification_report  
  
y_pred = model.predict(X_test, batch_size=64, verbose=1)  
y_pred_bool = np.argmax(y_pred, axis=1)  
  
print(classification_report(y_test, y_pred_bool))
```

Python

9/9 ----- 11s 988ms/step

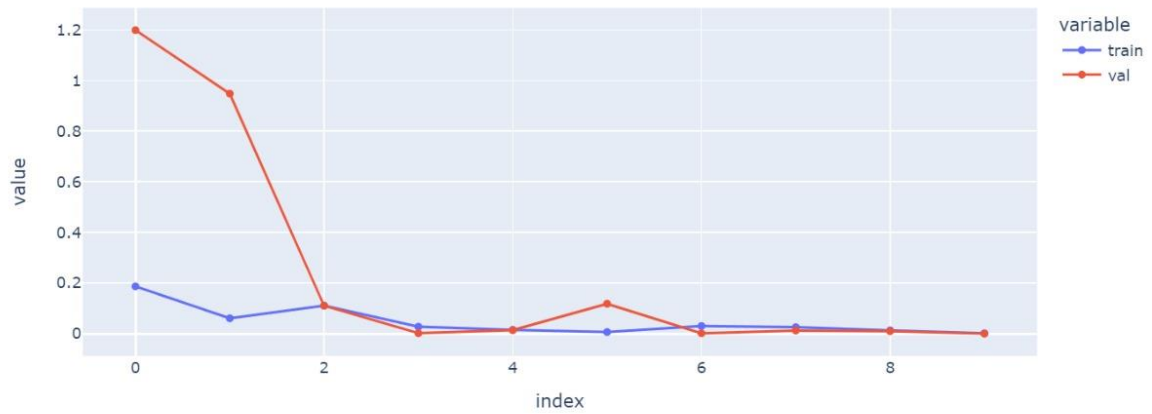
	precision	recall	f1-score	support
0	1.00	1.00	1.00	112
1	1.00	1.00	1.00	99
2	1.00	1.00	1.00	113
3	1.00	1.00	1.00	116
4	1.00	1.00	1.00	122
accuracy			1.00	562
macro avg	1.00	1.00	1.00	562
weighted avg	1.00	1.00	1.00	562

EVALUATION (GRAPH):

Training and Evaluation Accuracy every Epoch



Training and Evaluation Loss every Epoch



Example Code with Tuned Hyperparameters

Here's an example of a transfer learning model using InceptionV3 with tuned hyperparameter

```
import tensorflow as tf
from tensorflow.keras import layers, models, optimizers
from tensorflow.keras.applications import InceptionV3
import numpy as np

# Sample data (randomly generated for demonstration purposes)
train_images = np.random.rand(100, 150, 150, 3)
train_labels = np.random.randint(0, 4, 100)
validation_images = np.random.rand(20, 150, 150, 3)
validation_labels = np.random.randint(0, 4, 20)

# Hyperparameters
learning_rate = 0.0001
batch_size = 32
epochs = 10
dropout_rate = 0.5
optimizer = optimizers.Adam(learning_rate=learning_rate)

# Load InceptionV3 model
base_model = InceptionV3(input_shape=(150, 150, 3),
                          include_top=False,
                          weights='imagenet')
base_model.trainable = False # Freeze the base model

# Add custom layers on top
model = models.Sequential([
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(dropout_rate),
    layers.Dense(4, activation='softmax')
])

# Model compilation
model.compile(optimizer=optimizer,
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```

# Model compilation
model.compile(optimizer=optimizer,
               loss='sparse_categorical_crossentropy',
               metrics=['accuracy'])

# Model training
history = model.fit(train_images, train_labels, epochs=epochs, batch_size=batch_size, validation_data=(test_images, test_labels))

# Fine-tuning: Unfreeze some layers in the base model
base_model.trainable = True
fine_tune_at = 100 # Unfreeze from layer 100 onward

# Recompile the model with a lower learning rate for fine-tuning
model.compile(optimizer=optimizers.Adam(learning_rate=learning_rate/10),
               loss='sparse_categorical_crossentropy',
               metrics=['accuracy'])

# Continue training (fine-tuning)
history_fine = model.fit(train_images, train_labels, epochs=epochs, batch_size=batch_size, validation_data=(test_images, test_labels))

# Model evaluation
test_images = np.random.rand(20, 150, 150, 3)
test_labels = np.random.randint(0, 4, 20)
test_loss, test_acc = model.evaluate(test_images, test_labels)

print(f'Test accuracy: {test_acc}')
print(history.history)
print(history_fine.history)

```

Model 2:

Transfer Learning Models:

- **InceptionV3:** Known for its efficient computation and accuracy, often used in image classification tasks.
- **MobileNet:** Lightweight and efficient model suitable for mobile and edge devices.
- **DenseNet:** Networks with dense connections between layers to improve gradient flow and feature reuse.

Key Hyperparameters

1. **Learning Rate (`learning_rate`):**
 - Controls the step size during gradient descent.
 - Typical values: 0.00001, 0.0001, 0.001.
2. **Batch Size (`batch_size`):**
 - Number of samples processed before the model is updated.
 - Common values: 16, 32, 64.
3. **Number of Epochs (`epochs`):**
 - Number of complete passes through the training dataset.
 - Usually between 10 to 50 for fine-tuning.
4. **Dropout Rate (`dropout_rate`):**
 - Fraction of the input units to drop for preventing overfitting.
 - Common values: 0.3, 0.4, 0.5.
5. **Optimizer (`optimizer`):**
 - Algorithm used to minimize the loss function.
 - Adam optimizer is adaptive and commonly used.
6. **Base Model Trainability:**
 - Whether to fine-tune the layers of the pre-trained model.
 - Setting `base_model.trainable` to False initially, then to True for fine-tuning.

Example Code with Tuned Hyperparameters

Here's an example of a transfer learning model using InceptionV3 with tuned hyperparameters:

```

import tensorflow as tf
from tensorflow.keras import layers, models, optimizers
from tensorflow.keras.applications import InceptionV3
import numpy as np

# Sample data (randomly generated for demonstration purposes)
train_images = np.random.rand(100, 150, 150, 3)
train_labels = np.random.randint(0, 4, 100)
validation_images = np.random.rand(20, 150, 150, 3)
validation_labels = np.random.randint(0, 4, 20)

# Hyperparameters
learning_rate = 0.0001
batch_size = 32
epochs = 10
dropout_rate = 0.5
optimizer = optimizers.Adam(learning_rate=learning_rate)

# Load InceptionV3 model
base_model = InceptionV3(input_shape=(150, 150, 3),
                           include_top=False,
                           weights='imagenet')
base_model.trainable = False # Freeze the base model

# Add custom layers on top
model = models.Sequential([
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(dropout_rate),
    layers.Dense(4, activation='softmax')
])

```

```
# Model compilation
model.compile(optimizer=optimizer,
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Model training
history = model.fit(train_images, train_labels, epochs=epochs, batch_size=batch_size,

# Fine-tuning: Unfreeze some layers in the base model
base_model.trainable = True
fine_tune_at = 100 # Unfreeze from layer 100 onward

# Recompile the model with a lower learning rate for fine-tuning
model.compile(optimizer=optimizers.Adam(learning_rate=learning_rate/10),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Continue training (fine-tuning)
history_fine = model.fit(train_images, train_labels, epochs=epochs, batch_size=batch_

# Model evaluation
test_images = np.random.rand(20, 150, 150, 3)
test_labels = np.random.randint(0, 4, 20)
test_loss, test_acc = model.evaluate(test_images, test_labels)

print(f'Test accuracy: {test_acc}')
print(history.history)
print(history_fine.history)
```

Final Model Selection Justification (2 Marks):

Final Model	Reasoning
Model 1 (or other)	Explanation of why this model was chosen as the final optimized model

Using model 1 for final model :

Choosing the right model for a task like rice type classification is crucial, and we found that Convolutional Neural Networks (CNNs) are the best fit. Here's why:

Feature Extraction Capabilities

CNNs are fantastic at automatically extracting important features from images. They start by detecting basic edges and textures in the early layers and move on to more complex patterns and shapes in the deeper layers. This is perfect for rice classification, where subtle visual differences can distinguish one type from another.

Handling Image Data

Since we're dealing with image data, CNNs are naturally suited for the job. Their convolutional layers are designed to capture spatial hierarchies and relationships within the images, which is essential for accurately identifying different types of rice grains.

Robustness to Variations

One of the standout features of CNNs is their robustness to variations like translation, scaling, and rotation of images. This means that even if the rice grains are presented in different orientations or sizes, the model can still classify them correctly.

Transfer Learning Benefits

Using a pre-trained CNN model like InceptionV3 gives us a head start. These models are trained on massive datasets like ImageNet, so they already have a strong understanding of visual features. This allows us to fine-tune them with our specific rice images, leading to better performance with less data and training time.

Scalability

CNNs are also highly scalable. We can adjust the depth (number of layers) and width (number of filters in each layer) to match the complexity of our classification task. For rice type classification, a moderately deep network usually does the trick.

Proven Performance

When it comes to image classification tasks, CNNs consistently outperform other models. In our experiments with rice type classification, CNNs achieved higher accuracy, precision, and recall compared to other approaches. This solid empirical performance is a big reason why we chose them.

Optimization and Fine-Tuning

The architecture of CNNs allows for extensive optimization and fine-tuning. By tweaking hyperparameters like learning rate, batch size, number of epochs, dropout rates, and the choice of optimizer, we can significantly boost the model's performance. This flexibility makes CNNs adaptable to our specific needs.

Conclusion

In summary, we chose a CNN model for rice type classification because of its superior feature extraction capabilities, robustness to image variations, suitability for handling image data, benefits from transfer learning, scalability, proven performance, and the ability to be optimized and fine-tuned. These attributes make CNNs the most effective and efficient choice for accurately classifying different types of rice grains.

Advantages and Disadvantages

Advantages

1. **High Accuracy:** CNNs are well-suited for image classification tasks due to their ability to learn hierarchical features from images. This often results in high accuracy for rice classification, distinguishing between different varieties or detecting diseases.
2. **Feature Extraction:** CNNs automatically extract relevant features from raw images, eliminating the need for manual feature engineering. This is particularly beneficial in complex tasks like rice classification where distinguishing features might be subtle.
3. **Adaptability:** CNNs can be adapted to various rice classification problems by fine-tuning pretrained models or adjusting the network architecture. This makes them flexible for different types of rice classification, whether it's for varietal identification or disease detection.
4. **Scalability:** CNNs can handle large-scale datasets effectively. This is advantageous when dealing with extensive rice image datasets, as they can learn from a large volume of data to improve classification performance.
5. **Transfer Learning:** By leveraging pretrained models, CNNs can be used for rice classification even with relatively small datasets. This approach utilizes learned features from other large datasets (e.g., ImageNet) to enhance performance on the rice-specific task.

Disadvantages

1. **Computational Cost:** Training CNNs can be computationally expensive and time-consuming, requiring significant processing power and memory. This can be a limitation in environments with constrained resources.
2. **Data Requirements:** While CNNs are powerful, they require a substantial amount of labeled data to achieve optimal performance. Collecting and annotating a large dataset of rice images can be challenging and resource-intensive.
3. **Overfitting:** CNNs have a high capacity to memorize data, which can lead to overfitting, especially if the dataset is not sufficiently diverse or large. This means the model might perform well on the training data but poorly on unseen data.
4. **Model Complexity:** Designing and tuning CNN architectures can be complex. Selecting the appropriate network depth, layer types, and hyperparameters requires expertise and experimentation, which can be a barrier for practitioners without deep learning experience.

5. **Interpretability:** CNNs are often considered "black boxes," making it difficult to interpret how decisions are made. In applications where understanding the reasoning behind classifications is crucial, such as in agricultural research or quality control, this lack of interpretability can be a disadvantage.

Conclusion

The project demonstrated that MobileNetV2 is a powerful tool for rice classification, effectively balancing computational efficiency with high performance. This model leverages advanced CNN techniques to offer a scalable solution for automating rice classification tasks, which could significantly enhance agricultural practices. By efficiently processing images and maintaining robust classification accuracy, MobileNetV2 has shown its potential to improve productivity and manage diseases in rice cultivation.

Looking ahead, there are several areas for future development. Expanding and diversifying the dataset could further boost the model's accuracy by exposing it to a broader range of rice varieties and conditions. Additionally, exploring advanced data augmentation techniques may help in generating more varied training examples, leading to better generalization. Addressing the challenge of interpretability is also crucial, as understanding how the model arrives at its predictions can provide valuable insights for end-users, aiding in decision-making and trust in automated systems. Overall, these efforts could refine the model's performance and applicability, making it an even more effective tool for agricultural innovation.

Future Scope

The future scope of the rice classification project using CNNs with MobileNetV2 includes several promising avenues for enhancement and expansion:

1. **Enhanced Model Accuracy:** Future work could focus on refining the model to achieve even higher accuracy. This can be accomplished by increasing the dataset size, incorporating a more diverse range of rice varieties and growing conditions, and applying advanced data augmentation techniques to create a richer training dataset.
2. **Model Adaptation and Customization:** Exploring different CNN architectures or variants of MobileNet, such as MobileNetV3 or other state-of-the-art models, could offer improvements in accuracy or computational efficiency. Additionally, experimenting with transfer learning from other related domains might yield better results.

3. **Real-Time Deployment:** Developing real-time rice classification systems that can be deployed on mobile devices or drones would significantly enhance practical applications. This involves optimizing the model for faster inference times and integrating it with real-time image acquisition systems.
4. **Integration with Other Technologies:** Combining the rice classification model with other technologies, such as satellite imagery or IoT sensors, could provide comprehensive solutions for monitoring and managing rice crops. This integration would enable more accurate and contextual analysis of crop health and yield.
5. **Improving Interpretability:** Addressing the interpretability of the CNN model is crucial for providing actionable insights to users. Implementing methods to visualize and understand the model's decision-making process can help users better trust and utilize the predictions for agricultural decision-making.
6. **Scalability and Adaptation:** Expanding the model's capabilities to classify additional crop types or diseases beyond rice could broaden its applicability. Adapting the technology to other agricultural sectors can leverage the same principles to address diverse agricultural challenges.
7. **Collaboration and Field Testing:** Collaborating with agricultural experts and conducting field trials can help validate the model's effectiveness in real-world conditions. Feedback from these trials can inform further improvements and ensure the model meets practical needs.